

# Problem and Data

## Anomaly Detection on Taxi Calls

We are contacted by a Taxi company:



## Anomaly Detection on Taxi Calls

**They have historical data about taxi calls in NYC**

- In particular, they recorded the number of calls
- ...Over regular time intervals

**A major decision for the company is choosing the size of the car pool**

- This depends on how many calls are expected
- ...So, we'd like to figure that out

**Moreover, sometimes the number of calls deviates from the usual patterns**

- The company is interested in detecting such "anomalies"
- ...And anticipating them, if possible

**We will focus on the task of detecting anomalies**

How would you start to tackle this problem?

## Getting Started

## A couple of good ideas:

*Talking to the customer* to understand:

- Their priorities
- How their business works
- Any expectation on the data
- ...

...And also *inspecting the data*

- ...So that we get a "feel" of how it works
- Formally: until we understand better its *statistical distribution*

**Doing both these things *early* is always a good idea**

## Basic Setup

**Let us start by setting up the notebook:**

```
In [1]: %load_ext autoreload
        %autoreload 2
```

Our module contains a pre-built function to load the data:

```
def load_series(file_name, data_folder):
    ...
```

- We will use data from the [Numenta Anomaly Benchmark \(NAB\)](#)
- NYC taxi data `nyc_taxi.csv` is in the `data/realKnownCause` folder

```
In [2]: from util import util # Import our submodule
        data_folder = '../data/nab'
        file_name = 'realKnownCause/nyc_taxi.csv'
        data, labels, windows = util.load_series(file_name, data_folder)
```

## A Look at the Data

**Let's have a look at all the data we loaded**

```
In [5]: data.head()
```

Out [5] :

	value
timestamp	
2014-07-01 00:00:00	10844.0
2014-07-01 00:30:00	8127.0
2014-07-01 01:00:00	6210.0
2014-07-01 01:30:00	4656.0
2014-07-01 02:00:00	3820.0

- `data` is a pandas `DataFrame` object
- It is essentially a *table*, in this case representing a *time series*
- There are well defined *column names* (here "value")
- There is a well defined row *index* (here "timestamp")
- Jupyter displays DataFrame objects as HTML tables

## Time Series and Pandas

### Our data is a *time series*

I.e. a sequence whose index represents time

- Specifically, we have a *univariate* time series...
- ...Since we are tracking only quantity (i.e. one variable)

### Times series have one difference w.r.t. classical table datasets

- ...I.e. their row index is *meaningful*
- Since it represents the *position* of the example in the sequence

### That said, we do not care about *how* time is represented

- Hence, time series are stored *just as usual*!
- Their peculiarities arise when we start to manipulate them

## Time Series and Pandas

### In pandas:

- Time series are stored as usual, via `DataFrame` or `Series` objects
- ...You just need to pay more attention to the index

It *may* be convenient using a `datetime` index

- A `datetime` object in python allows to manipulate dates/hours directly
  - E.g. get year/month/day/hour/minute...
- In pandas they can be used as indices, so that for example:
  - Time stamps are easier to read
  - We can sort rows by time
  - We can represent arbitrarily long gaps between measurements
  - ...


That said, we still deal with normal `DataFrame` or `Series` objects

## A Look at the Data

**Let's have a look at all the data we loaded**

Our module contains a function to plot NAB series:

```
In [6]: util.plot_series(data)
```



## A Look at the Data

**Let's have a look at all the data we loaded**

We can now move to other data structures

```
In [7]: labels.head()
```

```
Out[7]: 0    2014-11-01 19:00:00
        1    2014-11-27 15:30:00
        2    2014-12-25 15:00:00
        3    2015-01-01 01:00:00
        4    2015-01-27 00:00:00
        dtype: datetime64[ns]
```

**labels** is a pandas `Series` object

- Similar to a 1D array
- ...But with a well defined row *index*

This series contains the *timestamp of all known anomalies*


- They are all hand-labeled

## A Look at the Data

Let's have a look at all the data we loaded

Let's plot both the series and the labels:

```
In [8]: util.plot_series(data, labels)
```



- Anomalies occur rarely (which is typical for this kind of problem)

## A Look at the Data

Let's have a look at all the data we loaded

Now the "windows" data structure:

```
In [9]: windows.head()
```

```
Out[9]:
```

	begin	end
0	2014-10-30 15:30:00	2014-11-03 22:30:00
1	2014-11-25 12:00:00	2014-11-29 19:00:00
2	2014-12-23 11:30:00	2014-12-27 18:30:00
3	2014-12-29 21:30:00	2015-01-03 04:30:00
4	2015-01-24 20:30:00	2015-01-29 03:30:00

**windows** is a pandas **DataFrame** object

- Contains the start/end of windows containing anomalies


- They represent a suitable "resolution" for detecting anomalies
- Reporting the presence of anomalies at any point of the window...
- ...Has some value for the company

## A Look at the Data

### Let's have a look at all the data we loaded

Let's plot the series, the labels, and the windows all together:

```
In [10]: util.plot_series(data, labels, windows)
```



- Detections that occur too early/late count as misses