```
# Notebook setup: run this before everything
        %load ext autoreload
        %autoreload 2
        # Control figure size
        figsize=(14, 4)
        from sklearn.neighbors import KernelDensity
        from sklearn.model selection import GridSearchCV
        from util import util
        import numpy as np
        from matplotlib import pyplot as plt
        import pandas as pd
        import os
        # Load data
        data_folder = os.path.join('..', 'data', 'nab')
        file_name = os.path.join('realKnownCause', 'nyc_taxi.csv')
        data, labels, windows = util.load series(file name, data folder)
        # Train and validation end
        train end = pd.to datetime('2014-10-24 00:00:00')
        val_end = pd.to_datetime('2014-12-10 00:00:00')
        # Cost model parameters
        c_alrm = 1 # Cost of investigating a false alarm
        c_missed = 10 # Cost of missing an anomaly
        c late = 5 # Cost for late detection
        # Build a cost model
        cmodel = util.ADSimpleCostModel(c alrm, c missed, c late)
        # Compute the maximum over the training set
        trmax = data[data.index < train_end]['value'].max()</pre>
        # Normalize
        data['value'] = data['value'] / trmax
        # Separate the training data
        data_tr = data[data.index < train_end]</pre>
        # Apply a sliding window
        wdata = util.sliding window 1D(data, wlen=10)
```

Time Indexed Models

Exploiting Time

Let's consider how we dealt with time so far

- We learned an estimator for f(t, x) and one for f(t)
- ...Which we used to compute $f(x \mid t) = f(t, x)/f(t)$

It worked well, but we had to introduce one additional dimension

What if we wanted to consider time and sequence input?

Let's consider a second approach to handle time

- This consists in *learning many density estimators*:
- Each estimator is specialized for a given time (e.g. 00:00, 00:30, 01:00...)

We can then choose which estimator to use based on the current time

Exploiting Time

Formally, what we have is a first ensemble model

In particular, we obtain our estimated probabilities by evaluating:

$$f_{g(t)}(x)$$

- Each f_i function is an estimator
- ullet The g(t) retrieves the correct f_i based (in our case) on the time value

We'll call this general idea a "selection ensemble"

In terms of properties:

- Each f_i estimator works with smaller amounts of data
- ...But the individual problems are easier!

Learning an Estimator for one Time Value

Let us make a test by learning an estimator for a single time value

First, we separate the training data

```
In [2]: wdata_tr = wdata[wdata.index < train_end]
wdata_tr.head()</pre>
```

Out[2]:	0	1	2	3	4	5	6
---------	---	---	---	---	---	---	---

timestamp								
2014-07- 01 04:30:00	0.357028	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	(
2014-07- 01 05:00:00	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	
2014-07- 01 05:30:00	0.204458	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	
2014-07- 01 06:00:00	0.153294	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	(
2014-07- 01 06:30:00	0.125770	0.094591	0.077997	0.067955	0.073124	0.071050	0.082804	(

- We'll use the normalized version
- ...So as to simplify our guesses for bandwidth selection

Learning an Estimator for one Time Value

Let us make a test by learning an estimator for a single time value

Then, we focus on the values for a single time value

```
In [3]: wdata_tr_test = wdata_tr.iloc[0::48] # 48 is the step
wdata_tr_test.head()
```

timestamp								
2014-07- 01 04:30:00	0.357028	0.267573	0.204458	0.153294	0.125770	0.094591	0.077997	C
2014-07- 02 04:30:00	0.440194	0.327429	0.249267	0.194811	0.158694	0.119646	0.098541	С
2014-07- 03 04:30:00	0.416357	0.347743	0.277088	0.233694	0.191815	0.144306	0.107661	C
2014-07- 04 04:30:00	0.513318	0.473941	0.412702	0.373391	0.328581	0.276693	0.237053	(
2014-07- 05 04:30:00	0.578672	0.533006	0.475455	0.412702	0.362361	0.301287	0.263721	С

2

3

5

Learning a 23:30 Estimator

0

Then we proceed as usual

We choose a bandwidth:

Out[3]:

- For sake of simplicity, we'll use the same bandwidth for all estimators
- ullet Even if we should re-calibrate h for each estimator in principle

Learning the Ensemble

Now, we need to repeat the process for every unique time value

```
In [6]: day_hours = data_tr.index.hour + data_tr.index.minute / 60
    day_hours = day_hours.unique()
    print(day_hours)
```

- unique in pandas returns a Series with all unique values
- We do not care about how time is measured
- ...We only care about having 48 discrete steps

Learning the Ensemble

Finally, we can learn 48 specialized estimators

```
In [7]: kde = {}
for hidx, hour in enumerate(day_hours):
    tmp_data = wdata_tr.iloc[hidx::48]
    kde[hour] = KernelDensity(kernel='gaussian', bandwidth=h)
    kde[hour].fit(tmp_data)
```

- For each unique time value, we separate a subset of the *training* data
- Then we build and learn a KDE estimator

We chose to store everything in a dictionary:

```
In [8]: print(str(kde)[:256], '...}')
```

 $\{0.0: KernelDensity(bandwidth=np.float64(0.019473684210526317)), 0.5: KernelDensity(bandwidth=np.float64(0.019473684210526317)), 1.0: KernelDensity(bandwidth=np.float64(0.019473684210526317)), 1.5: KernelDensity(bandwidth=np.float64(0.019473684210526317)), ...\}$

Generating the Signal

The we can generate the alarm signal

- In a practical implementation we should do this step by step
- ...But for an evaluation purpose it is easier to do it all at once

- For each unique time value, we separate a subset of the whole data
- Then we obtain the estimated (log) probabilities

The process is even faster than before

• ...Because each KDE estimator is trained a smaller dataset

Generating the Signal

All signals are stored in a list

- We need to concatenate them all in single DataFrame
- Then we can sort all rows by timestamp (it's the index)

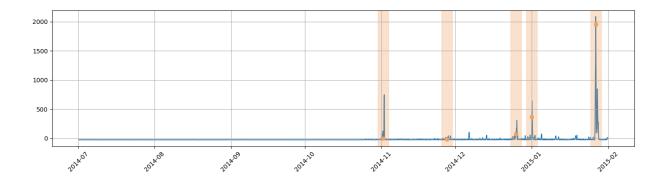
A suggestion: always do concatenations in a single step in pandas

It's way faster than appending DataFrame objects one by one

Generating the Signal

Now we can plot out signal:

```
In [11]: util.plot_series(signal, labels=labels, windows=windows, figsize=figsize)
```



- It's very similar to that of the other time-based model
- ...But also a bit smoother, like that of the sequence-based model

Threshold Optimization and Evaluation

Now we can optimize the threshold and evaluate the results

```
In [12]: signal_opt = signal[signal.index < val_end]
    labels_opt = labels[labels < val_end]
    windows_opt = windows[windows['end'] < val_end]
    thr_range = np.linspace(10, 200, 100)

best_thr, best_cost = util.opt_thr(signal_opt, labels_opt, windows_opt, cmoc print(f'Best threshold: {best_thr}, corresponding cost: {best_cost}')</pre>
```

Best threshold: 104.04040404040404, corresponding cost: 10

Let us see the cost on the whole dataset:

```
In [13]: ctst = cmodel.cost(signal, labels, windows, best_thr)
    print(f'Cost on the whole dataset {ctst}')
```

Cost on the whole dataset 10

This is the best result we have achieved so far!

What if we used this approach for the second period?